

Contract & SLA Copilot for Service Operations

An internal AI copilot that helps service teams quickly verify contract coverage and SLA commitments using cited answers from contract documents.

Prioritisation Score: 21/25 | 2x2:  Strategic Bet | Technique: Agent+RAG | Function: Legal / Compliance

PRIMARY USERS

Role	Department	Frequency of Use
Service Desk Lead / Contract Manager	Customer Service / Commercial	Daily - multiple queries per shift
Account Manager / Service Engineer	Sales / Field Service	Weekly - SLA reviews & escalations
Legal / Commercial Analyst	Legal / Finance	Ad-hoc - contract audits & renewals
These roles collectively generate a steady stream of contract interpretation queries across service operations.		

DECISION / WORKFLOW IMPROVED

Answering customer contract and SLA queries - specifically decisions on coverage eligibility, applicable SLA response times, and credit/penalty entitlements. Currently a manual search through lengthy service agreements; replaced by an instant cited draft answer that the employee can verify before responding to the customer.

PROBLEM STATEMENT & SOLUTION

Service desk teams handling customer contract queries currently spend 10–30 minutes per inquiry manually searching through complex, multi-page service agreements and SLA tables to answer questions like "Is this failure mode covered?", "What is our committed response time?" or "What penalty credits apply?" This process is slow, inconsistent across agents, and frequently results in wrong commitments - creating legal and financial exposure. As contract portfolios scale, the manual lookup burden becomes a direct bottleneck to customer satisfaction, commercial accuracy, and operational efficiency. **The solution** uses an Agent-based Retrieval-Augmented Generation (RAG) architecture, where specialized agents handle query understanding, document retrieval, reasoning over contract clauses, and escalation to human review when uncertainty is detected.

WHY GenAI (Agent+RAG Approach)

- **Contextual Reasoning over Complexity:** Service contracts are written in natural language with complex cross-references. Our **Reasoning Agent** evaluates these documents to **understand how multiple clauses interact**, resolving nuances and ambiguities that rule-based systems would miss.
- **Semantic Intent Matching:** Queries arrive as unstructured natural language. RAG matches a user's underlying intent to relevant contract sections-not just keywords-enabling more nuanced retrieval than classic ML classifiers.
- **Dynamic Knowledge without Retraining:** Contract repositories are in constant flux. RAG retrieves from the latest indexed documents instantly, eliminating the need for the expensive re-labeling and re-training cycles required by traditional models.
- **Safety via Agentic Escalation:** The workflow includes an Escalation Agent that detects ambiguity or missing data. If "legal ground truth" cannot be established, the system flags the case for human review, preventing incorrect commercial commitments.

KEY CONSTRAINTS

Governance (HITL)	Mandatory Human-in-the-Loop; GenAI provides drafted answers with citations, but all customer-facing commitments require human approval.
Data Classification	All contracts are Confidential . Data must remain within the approved cloud tenant; no external LLM processing without a valid Data Processing Agreement (DPA).
Scope Discipline	Limited to 3 question types: (1) Coverage, (2) SLA commitments, (3) Service credit eligibility. Non-features include multi-language support and auto-sending responses
Timeline & Approvals	4 week MVP target. Requires Legal sign-off on (a) 'Not Legal Advice' disclaimer, (b) escalation paths for ambiguity, and (c) evaluation set accuracy.

DEPENDENCIES

- **System access:** Read access to contract repository (SharePoint / DocuSign / DMS); no write access needed at MVP stage
- **Document readiness:** Contracts must be clean, text-selectable PDFs; scanned documents require OCR preprocessing before ingestion
- **LLM / RAG stack:** LLM API with enterprise data agreement (e.g., Azure OpenAI, Anthropic Claude API on-prem); RAG framework (LlamaIndex or LangChain); vector database (Pinecone, ChromaDB, or pgvector)
- **Identity & RBAC:** Integration with the existing identity provider (e.g., Azure AD / Okta) to enforce Role-Based Access Control, ensuring agents only retrieve contracts within their authorized scope.
- **Stakeholder Validation:** Appointment of Contract Managers to provide a "Golden Set" of 50–100 labeled test cases for accuracy benchmarking and Legal sign-off on the mandatory disclaimer language

EXPECTED VALUE (quantified range)

Value Driver	Low Scenario (Conservative)	Likely Scenario (Target)	High Scenario (Optimistic)
Annual Query Volume	2,000 queries	2,600 queries (≈50/week)	3,000 queries
Minutes Saved / Query	31 min	31 min	31 min
Annual Hours Saved	1,033 hours	1,343 hours	1,550 hours
Loaded Hourly Cost	€45	€55	€65
Time Savings Value	€46,500	€73,865	€100,750
Risk / Penalty Avoidance	Avoid 1 incident/year	Avoid 3 incidents/year	Avoid 5 incidents/year
Quality Savings (€)	€20,000	€60,000	€150,000
Total Annual Value	€66,500	€133,865	€250,750

 Total estimated annual value: €66,500 – €250,750+ (conservative driven by time savings alone; upside driven by risk/penalty avoidance)

PRIORITISATION SCORE

Value Potential	Time-to-Value	Feasibility	Risk/Constraints	Evaluability	TOTAL
4 / 5	4 / 5	4 / 5	4 / 5	5 / 5	21 / 25

This is ranked highest because the use case addresses a frequent real operational task involving large document collections and understanding. It has clear business value, strong evaluability because answers can be verified against contract clauses, and it demonstrates how **Agent AI with RAG** can support real enterprise workflows.

NEXT STEPS (before MVP design - Block B/C)

- **Audit Repository & Volume:** Conduct a systematic audit of the contract portfolio to verify document quality (text-selectable vs. scanned) and confirm weekly query volumes to ensure the "High Volume" filter is met.
- **Establish "Golden Set":** Compile an evaluation dataset of 50–100 real-world historical contract questions and expert-verified answers to serve as the ground truth for benchmarking. The MVP's success will be measured against this set by achieving **85% answer accuracy, 90% citation grounding, 100% escalation of ambiguity**, sub-60-second response times, and a **50% reduction in manual lookup time**.
- **Define Legal & Compliance Scope:** Secure formal sign-off on the specific contract types in scope, mandatory "Not Legal Advice" disclaimers, and defined escalation paths for high-risk or ambiguous clauses.
- **Run a Feasibility Pilot:** Execute a rapid RAG test on 5–10 representative contracts to verify the architecture can achieve a baseline retrieval accuracy of ≥85%.
- **Validate User Friction:** Interview 3–5 service desk agents to map top frustration points and verify baseline handle times required for the Block B Value Sizing.

Block B: B1–B4 | Process Mapping and Value Sizing

B1 — Process Map (As-Is / Current Process)

Process Boundary:

Start Trigger: Customer or internal agent submits a query regarding SLA commitments, contract coverage, or service credit eligibility.

End State: A final response is delivered to the customer and the interaction is logged in the service management system.

Total Baseline Time: Exactly 37 minutes

Step	Role	Input	Output	Avg Time	Pain / Failure	Risk	Source / Justification
1	Service Agent	Customer question (email/ticket/phone)	Clarified query description	4 min	Query often vague; agent may misinterpret the issue	Wrong problem framing leads to wrong clause search	Typical service desk triage behavior in customer support workflows
2	Service Agent	Customer ID / ticket / CRM record	Contract ID and service context	5 min	Information scattered across CRM, ticketing, asset database	Wrong contract selected or missing entitlement context	Enterprise service desks commonly require manual account lookup
3	Service Agent	Contract repository (SharePoint / DMS)	Correct contract document and SLA appendix	6 min	Version confusion; amendments stored in different folders	Outdated contract version used	Document management systems often contain multiple contract versions
4	Service Agent	Contract PDFs and SLA tables	Potential relevant clauses	8 min	Manual reading of long documents	Clause overlooked or misunderstood	Legal documents require interpretation and cross-references
5	Contract Manager / Specialist	Clauses identified by agent	Interpreted coverage / SLA rule	6 min	Inconsistent interpretation across experts	Incorrect legal/commercial commitment	Contract interpretation often requires specialist knowledge
6	Service Agent	Specialist guidance and clause text	Draft response to customer	4 min	Draft may lack full context or precise wording	Customer receives unclear or incorrect explanation	Manual drafting introduces variability
7	Manager / Senior Reviewer	Draft answer	Approved final response	3 min	Late-stage review causes rework	Risk exposure if review is skipped	Escalation practices common in commercial service operations
8	Service Agent	Approved answer	Customer reply + ticket update	1 min	Logging inconsistent or incomplete	Poor auditability and weak knowledge capture	Manual CRM/ticket logging typical in service workflows

The current process relies heavily on manual document retrieval and contract interpretation, creating delays and increasing the risk of inconsistent commitments. The primary bottlenecks occur during document search and clause interpretation (Steps 3–5), which together account for over half the total resolution time. The lack of structured retrieval and standardized interpretation increases both response time and the risk of incorrect commitments.

B2 — Process Map (To-Be with GenAI)

Process Boundary:

Start Trigger: Customer or agent submits a contract or SLA question in the service system.

End State: A verified, cited answer is sent to the customer and automatically logged for monitoring and audit.

Target Time ≈ 6 minutes

Step	Role	Avg Time (In min.)	AI Step? (Y/N)	Human Approval?	What gets logged	Output
1	Service Agent	0.5	N	N	Query timestamp, ticket ID, user ID	Customer SLA / contract question submitted
2	Extraction Agent	0.5	Y	N	Query classification, extracted entities (customer, contract ID, asset)	Structured query summary
3	Retrieval + Reasoning Agent	1.0	Y	N	Retrieved documents, cited clauses, confidence score	Relevant SLA clauses and contract sections
4	Reasoning Agent	1.0	Y	N	Draft answer, reasoning trace, ambiguity flag	Customer-ready drafted response with citations
5	Service Agent	2.0	N	Y	Reviewer ID, edits, acceptance decision	Verified response ready for sending
6	Contract Manager (only if flagged)	0.0	N	Y	Escalation reason, interpretation note	Specialist-approved clarification
7	Service Agent	0.5	N	Y	Final response text, send timestamp	Final answer delivered to customer
8	System Logging Layer	0.5	Y	N	Full interaction log (query, clauses used, draft, edits, final answer)	Case stored for monitoring, QA, and evaluation dataset

The redesigned process introduces GenAI agents for query structuring, document retrieval, and draft generation, significantly reducing manual lookup and interpretation time. Human oversight remains in the review and response stages, preserving accountability and preventing incorrect commitments. Automated logging creates a structured audit trail that supports governance, quality monitoring, and future evaluation datasets.

Figure 1: Current (“As-Is”) contract and SLA query handling process.

Customer contract or SLA questions are currently handled through a largely manual workflow involving document retrieval, clause interpretation, and internal review. Service agents search multiple systems and contract repositories to locate relevant clauses, often requiring specialist interpretation before drafting and approving a response. This sequential process creates delays and increases the risk of inconsistent interpretation or incorrect commitments.





Figure 2: Future (“To-Be”) workflow with GenAI-supported contract and SLA interpretation.

The redesigned workflow introduces GenAI agents for query extraction, document retrieval, and draft response generation with clause-level citations. Human oversight remains in the review and response stages to ensure accountability and prevent incorrect commitments. Automated logging captures queries, citations, edits, and final responses, creating an auditable knowledge trail while reducing resolution time from approximately 37 minutes to about 6 minutes per query.



B3 — Value Sizing Table (The value sizing model below provides the transparent calculation behind the annual value range summarized in Block A)

Assumption / Calculation Step	Low Scenario	Likely Scenario	High Scenario	Source / Justification
Annual Query Volume	2,000 queries	2,600 queries	3,000 queries	Baseline service-desk workload estimate; likely case anchored on ~50 queries/week
As-Is Handling Time	37 min	37 min	37 min	From validated As-Is process map
To-Be Handling Time	6 min	6 min	6 min	From validated To-Be process map
Minutes Saved / Query	31 min	31 min	31 min	Calculated as 37 – 6
Annual Hours Saved	1,033 hours	1,343 hours	1,550 hours	Formula: Volume × 31 / 60
Loaded Hourly Cost	€45	€55	€65	Loaded labor cost range for service / contract operations roles
Annual Time Value	€46,500	€73,865	€100,750	Formula: Annual Hours Saved × Hourly Cost
Quality / Risk Savings	€20,000	€60,000	€150,000	Reduced wrong commitments, avoided SLA penalties, less rework and dispute handling
Total Annual Value	€66,500	€133,865	€250,750	Formula: Time Value + Quality / Risk Savings

Annual Time Value = (Volume x Minutes Saved/60) x Hourly Cost

Example (Likely scenario):(2600 × 31/60) × 55 = €73,865

This value reflects **productive time freed from manual contract lookup and interpretation.**

Value per Query: Each query resolved using the GenAI copilot saves approximately **31 minutes**, translating to **€23–€34 of operational value per query** depending on the loaded hourly cost.

Value per Query = (Minutes Saved/60) x Hourly Cost

Example (Likely scenario): (31/60) × €55 = €28.4

With baseline ~50 queries/week 50 queries/week × €28.4 value ≈ €1,400 value/week

B4 — 6-Line Business Case Narrative

Value is primarily driven by a **31-minute reduction in handle time per query** and the mitigation of financial risk by avoiding incorrect SLA commitments. This impact is measurable through automated logs tracking **response times, citation accuracy, and human override rates**, providing a transparent audit trail for ROI verification. Our primary assumption of **50 queries per week** will be validated by auditing historical service logs and monitoring real-time adoption during the pilot phase. To mitigate the risk of complex clause misinterpretation, we have implemented **mandatory human review and automated escalation** to contract specialists when ambiguity is detected. Success after 30 days is defined as a **greater than equal to 50% reduction in manual lookup time**, greater than equal to 85% answer accuracy, and high agent acceptance. We will stop or pivot the project if **citation accuracy falls below acceptable thresholds** or if real-world query volume is insufficient to justify the operational investment.

C1. MVP PRD v1 - Contract & SLA Copilot for Service Operations

An internal AI copilot that helps service teams quickly verify contract coverage and SLA commitments using cited answers from contract documents.

1.1 Product Identity

Product Name	Contract & SLA Copilot for Service Operations	
Target Users	Primary: Tier-1 / Tier-2 Customer Service Agents (EMEA service operations) (EMEA Tier-1/2 Service Agents are frontline customer support professionals responsible for handling initial service requests and SLA-related queries for customers in the Europe, Middle East, and Africa (EMEA) region. They perform first-level triage, interpret standard contract and SLA terms (e.g., coverage, response times, service credits), and escalate complex or ambiguous cases to contract managers or specialists.) Secondary: Contract Managers (escalation review), Client-facing Account Managers (B2B), Operations Managers needing SLA verification.	
Usage Context	Agents receive customer queries related to: - SLA commitments - service coverage - service credit eligibility - warranty entitlements	Agents currently must manually: - locate the correct contract - search SLA clauses - interpret legal wording - draft a response This takes ~37 minutes per query in the current process.
Core JTBD (Job-to-be-Done)	When a customer submits a support request about contract coverage or SLA commitments, I want to instantly see the relevant contract clauses and a suggested response draft with citations, so that I can verify the answer quickly and respond accurately without manually searching through contracts.	
Problem Statement	Customer service agents currently resolve contract-related queries through a manual process involving document retrieval, clause interpretation, and response drafting. This process takes approximately 37 minutes per query , with most time spent locating documents and interpreting SLA clauses.	
Solution	The Contract & SLA Copilot reduces this bottleneck by retrieving relevant contract clauses and generating a cited response draft , allowing agents to review and respond in approximately 6–10 minutes , consistent with the redesigned To-Be workflow.	

1.2 Why Generative AI?

GenAI is required because the problem involves **interpreting unstructured legal documents and producing natural-language explanations**.

Reasons:

- Contracts are unstructured documents:** SLA clauses, appendices, and amendments vary widely in format and wording.
- Queries require contextual reasoning:** Answers often depend on interpreting multiple clauses together.
- Agents need customer-ready explanations:** The system must produce a clear response draft supported by clause citations.
- Answers must be **grounded in verifiable clause citations**, which requires combining retrieval with natural-language reasoning.

1.3 In-Scope Features (MVP ~4 weeks)

#	Feature	Working and Purpose
F1	Query Structuring	The system converts a customer question into a structured query. Example extraction: contract ID / customer name, SLA topic, service type, relevant asset or entitlement. Purpose: speed up contract retrieval; reduces wrong document selection.
F2	Contract Clause Retrieval with Citations	The system retrieves relevant contract sections using semantic search. Output includes clause snippets, clause identifiers, document references. Purpose: removes manual document search; ensures answers are grounded in real clauses.
F3	AI Draft Response with Human Review	The system generates a draft response with clause citations. Agents review the suggestion in a panel and can accept, edit, reject. Purpose: eliminates manual drafting; keeps human accountability before sending.

1.4 Non-Features (MVP Boundary)

To keep the build feasible within **4 weeks**, the following features are explicitly excluded.

- Automatic sending of customer responses:** All outputs require manual review.
- Multi-language support**
MVP supports **English only**.
- Full CRM or ticketing system integration**
- Voice call transcription**
- Contract editing or lifecycle management**
- Automated SLA breach prediction**
- Mobile application**
- Advanced analytics dashboards**

These items are potential **Phase-2 enhancements** but are intentionally excluded to maintain scope discipline.

1.5 User Stories

ID	As a ...	I want to/the ...	so that ...
US-01	Tier-1 service agent	see relevant contract clauses automatically	I can verify SLA coverage quickly.
US-02	Tier-1 service agent	see an AI-generated response draft with clause citations	I can answer customer queries faster.
US-03	Service agent	edit the AI suggestion before sending	I remain accountable for the final message.
US-04	Contract specialist	system to flag ambiguous cases	complex interpretations can be reviewed.
US-05	Support manager	see whether suggestions were accepted or rejected	the pilot's effectiveness can be evaluated.

1.6 Output Format & Constraints

Output Format	The system displays results in a copilot suggestion panel next to the support ticket. The panel contains: Suggested response draft Cited contract clauses Clause references and document location Clause-grounding confidence indicator (flags low-confidence answers for specialist review) Accept / Edit / Reject actions
---------------	--

Constraints	Latency: Response generation must complete in ≤ 60 seconds Citation requirement: Every response must include explicit clause references Human-in-the-Loop: AI suggestions must be reviewed by an agent before sending Data scope: Retrieval limited to internal contract documents (Master Service Agreements, SLA appendices, Contract amendments, Warranty policies) Security: No external model training using contract data. Logging: The system logs query text, retrieved clauses, model output, agent edits, and final response for audit and evaluation.
-------------	---

1.7 Acceptance Criteria (Testable)

AC#	Criterion	Pass/Fail Signal
AC-01	Answer accuracy: Evaluation using a 50-query golden evaluation set derived from historical service tickets and validated by contract specialists	≥ 85% answers correctly interpret contract clauses
AC-02	Citation Validity	At least 90% of responses include correct clause citations
AC-03	Latency	95% of queries produce responses within 60 seconds
AC-04	Handle-Time Reduction: Average resolution time decreases from	37 minutes → ≤ 10 minutes during the pilot phase.
AC-05	Agent Adoption	At least 70% of AI suggestions are accepted or edited rather than rejected.


C2. Architecture Sketch v1

Product: Contract & SLA Copilot for Service Operations (A vendor-neutral Agentic Workflow architecture with Retrieval-Augmented Generation (RAG) as the knowledge-retrieval mechanism, enabling clause-level contract reasoning with human approval gates.)

Pattern 4: Agentic Workflow used because the system must: decompose the task into steps; call different components in sequence; decide whether ambiguity exists; pause for human approval before any high-impact action; and maintain auditability.

2.1 System Components

Component	Purpose & Functionality	Outputs & Impact
1. UI / Copilot Panel	Lightweight panel embedded beside the service ticket. Captures questions, displays AI drafts with citations, and enables Accept/Edit/Reject actions.	Output: Suggested answer, citations, and confidence/ambiguity flags. Impact: Eliminates manual drafting and searching.
2. Orchestration Layer	The central controller coordinating all steps. Triggers agents, applies routing logic, and sends flagged cases to human review.	Output: Governed data flow and final UI outputs. Impact: Provides structural integrity and ensures consistency.
3. Extraction Agent	A preprocessing agent that converts raw text into structured data (Customer ID, SLA Topic, Entitlement type).	Output: Structured query package and extracted entities. Impact: Reduces "wrong-document" selection and manual triage time.
4. Retrieval + Reasoning Agent (RAG Engine)	Performs semantic retrieval from the vector database and generates grounded answers using clause-level evidence.	Output: Grounded draft answer with clause-level citations. Impact: Automates the largest bottleneck: manual interpretation.
5. Vector Database	A retrieval store containing indexed contract/SLA chunks and metadata (version, ID, section).	Output: Filtered, searchable document embeddings. Impact: Enables fast semantic retrieval and full citation traceability.
6. Source Document Indexing Layer	A background service that ingests, chunks, and indexes approved contracts.	Output: Searchable knowledge base. Impact: Ensures the RAG system uses the latest "read-only" document versions.
7. System Logging Layer	Records the full workflow (inputs, outputs, user actions) for audit and evaluation.	Output: Audit trail and performance metrics (latency, accuracy). Impact: Supports the "Golden Set" and governance-lite needs.



2.2 Data Sources (MVP Read-Only Access)

Source	Examples	Purpose & Usage
1. Contract Repository	SharePoint, DMS, Approved Contract Folders	Accesses Master Service Agreements (MSAs), signed amendments, annexes, and version-controlled contract files.
2. SLA Entitlement Tables	SLA Schedules, Service Credit Tables, Warranty Lists	Provides committed response times, eligibility thresholds, and specific service credit rules for entitlement verification.
3. Ticket/Case Metadata	Ticket ID, Customer Identifier, Issue Category	Minimal metadata used to structure the query and route the request; kept limited to avoid excessive integration work during the 4-week MVP.

2.3 End-to-End Flow

Step	Phase	Action & Component Involvement	Output & Impact
1	Query Submission	Agent enters/pastes the customer's SLA or contract question into the UI Panel .	Output: Raw query capture.
2	Query Structuring	Orchestration Layer sends input to the Extraction Agent to identify intent and entities.	Output: Structured query summary.
3	Retrieval	Retrieval + Reasoning Agent searches the Vector Database for relevant clauses.	Output: Relevant contract/SLA clause set.
4	Draft Generation	Reasoning Agent drafts a response with citations and ambiguity/confidence flags.	Output: Cited draft with rationale trace.
5	Human Review	Service Agent reviews the answer in the UI Panel ; decides to Accept, Edit, or Reject .	Output: Verified response draft.
6	Specialist Review	If flagged for ambiguity or conflict, the case is routed to the Contract Manager .	Output: Specialist-approved clarification.
7	Final Response	The verified answer is delivered through the normal service workflow by the service agent.	Output: Final response to customer.
8	Logging	System Logging Layer stores the full interaction (input, draft, edits, and citations).	Output: Audit trail and Golden Set enrichment.

2.4 Where Humans Approve

Approval Gate	Location in Flow	Strategic Justification
1. Service Agent	Post-Draft Generation: After the Reasoning	Preserves human accountability and prevents unreviewed commercial

Review	Agent produces a cited response.	commitments. Directly satisfies the Human-in-the-Loop (HITL) requirement.
2. Specialist Review	Escalation Only: Triggered only when the agent flags ambiguity or clause conflicts.	Ensures high-risk legal interpretations are handled by experts. Allows escalation for cases where "legal ground truth" cannot be autonomously established.

2.5 What Gets Logged (Audit Trail)

Category	Data Fields to Capture	Strategic Purpose
Inputs	Raw user query, Ticket ID, User ID/Role, and Timestamp.	Establishes the initial audit trail and identifies the primary user for each case.
Intermediate System Data	Extracted entities, classified query type, retrieved document IDs, cited clause references, confidence/ambiguity flags, and step-by-step latency.	Allows for technical troubleshooting and identifies exactly where the "Reasoning Agent" may be struggling.
Outputs	Initial draft answer, retrieved clause snippets, and the final cited answer displayed to the agent.	Tracks the raw performance of the RAG system before human intervention.
Human Actions	Decision (Accepted/Edited/Rejected), specific text edits made by the agent, specialist escalation flags, and specialist override notes.	Measures the "Human-in-the-Loop" impact and identifies common areas where the AI draft requires correction.
Outcomes	Final resolution time, citation accuracy result, and estimated handle-time saved per query.	Directly supports the 50–100 case Golden Set and supports validation of the estimated €133,865 annual operational value identified in Block B.

2.6 Top 5 Failure Modes & Mitigations

Failure Mode	Identified Risk	Architectural Mitigation Strategy
1. Retrieval Misses Right Clause	System retrieves irrelevant sections or misses the controlling clause entirely.	Use Hybrid Retrieval (semantic + keyword) and reranking ; implement mandatory contract-version filtering to ensure the latest amendment is used.
2. Hallucinated / Overconfident Answer	Model drafts a plausible answer that is not supported by the source contract text.	Enforce grounding rules where the agent only answers from retrieved evidence; require mandatory citations and human review before any customer-facing action.
3. Ambiguous / Conflicting Clauses	Multiple clauses or amendments appear to conflict, leading to incorrect commitments.	Implement an ambiguity detection logic within the Reasoning Agent ; trigger an automated escalation path to a Contract Manager for manual specialist review.
4. Latency Exceeds 60 Seconds	High response times reduce agent usability and harm pilot adoption.	Pre-index documents offline ; use extraction agents to narrow the retrieval scope; separate the ingestion layer from the live inference path.
5. Unauthorized Access to Contracts	A user retrieves confidential contracts outside of their authorized scope or department.	Integrate Role-Based Access Control (RBAC) ; apply metadata filters at the retrieval step to restrict data access based on the logged user ID.

C3. Build / Buy Decision Table

The decision context hinges on the system's need for an **Agentic workflow** that integrates structured extraction, retrieval, and reasoning with an automated escalation path ; the architecture must ensure strict **RAG retrieval over confidential contract documents** while maintaining an explicit **Human-in-the-Loop approval** process for all customer-facing commitments ; functionality requires granular, **clause-level citations** to provide a verifiable "legal ground truth" ; performance is bound by a strict **60-second latency** target to ensure usability for the **50+ internal service agents** ; finally, the evaluation must determine if commodity off-the-shelf tools can handle these specialized legal reasoning requirements or if a custom-built solution is necessary to achieve the desired **85% answer accuracy**.

Criterion	Build	Buy	Notes
Speed to MVP (4-week constraint)	Medium – requires architecture setup, document indexing, and UI panel development	High – existing copilots and AI search tools can be deployed quickly	Buying accelerates initial deployment but may require configuration to match the workflow
Agentic Workflow Support	High – custom orchestration layer can fully implement the 8-step To-Be process	Medium – many vendor tools support RAG but limited multi-step reasoning workflows	Our architecture requires structured extraction, reasoning, and escalation logic
Contract Interpretation Accuracy	High – retrieval rules and reasoning prompts can be tuned specifically for SLA clauses	Medium – generic assistants may struggle with legal interpretation	Contract interpretation is a domain-specific reasoning problem
Citation Grounding Requirement	High – build allows strict enforcement of clause-level citations	Medium – some tools provide citations but not always clause-level traceability	Accurate citation is essential for compliance and customer commitments
Security of Confidential Contracts	High – internal deployment with strict RBAC controls	Medium – vendor systems require careful configuration to protect sensitive data	Contracts contain confidential commercial terms
Integration with Internal Document Repository	Medium – requires indexing pipeline development	High – many AI search platforms integrate directly with document systems	For MVP, minimal integration may be sufficient
Scalability for Future Automation	High – architecture can evolve to include SLA monitoring, analytics, and automation	Medium – vendor limitations may restrict future workflow extensions	Building provides long-term flexibility
Operational Maintenance Effort	Medium – requires internal support and monitoring	Low – vendor maintains infrastructure and updates	Maintenance cost must be considered
Cost vs. Annual Value	Medium – development cost upfront but lower recurring licensing	Medium–High – vendor licensing costs may accumulate	Value case must justify ongoing cost
Differentiation / Competitive Advantage	High – system becomes internal IP for contract intelligence	Low – competitors could use the same off-the-shelf tools	Legal interpretation capability could become strategic

Recommended Approach: Hybrid (Buy Core Infrastructure; Build the Intelligence Layer)

The **Contract & SLA Copilot** requires a balance between **commodity efficiency** and **strategic differentiation**. We recommend purchasing the underlying infrastructure—such as **document indexing, vector databases, and LLM inference**—to meet the strict **4-week MVP timeline**. However, the **Intelligence Layer**—including extraction logic for SLA queries, contract-specific reasoning prompts, and ambiguity detection—must be built internally. Generic copilots lack the clause-level precision and citation grounding required for legally sensitive service commitments.

Strategic Justification:

- Speed to Value:** Leveraging existing AI infrastructure enables deployment within the 4-week sprint.
- Security & Compliance:** Confidential contracts remain within internal environments with robust **RBAC and audit logging**.

- **Competitive Advantage:** The specialized reasoning layer becomes proprietary intellectual property, continuously improving accuracy and reducing risk.
- **ROI Alignment:** This path captures the **€133,865 annual value** identified in Block B while avoiding the prohibitive costs of building a full infrastructure stack

Estimated MVP Cost vs. Value Capture:

The expected annual value of the Contract & SLA Copilot is **€133,865**, primarily from reduced agent handle time and fewer incorrect SLA commitments (Block B).

The MVP implementation is intentionally lightweight and limited to core features:

- document ingestion and indexing
- vector retrieval infrastructure
- reasoning layer configuration
- internal copilot panel for ~50 agents

A typical internal MVP deployment of this architecture would require approximately:

Cost Element	Estimated Cost
Infrastructure setup	€8,000 – €12,000
Development effort (4 weeks small team)	€25,000 – €35,000
Pilot monitoring and evaluation	€5,000 – €8,000

Estimated MVP investment: €38,000 – €55,000.

Given the annual value of **€133,865**, the system would likely reach **payback within the first year of operation**, making the initiative economically attractive even before broader automation benefits are realized.

Strategic Risk Considerations:

Because the system processes **commercial contracts and SLA commitments**, the architecture must manage several operational risks.

Risk Area	Mitigation
Incorrect contract interpretation	Mandatory clause citations and human review before sending responses
Exposure of confidential contracts	Role-based access control and retrieval filters based on user permissions
Model hallucination	Retrieval-grounded answers only; responses generated strictly from retrieved clauses
Operational dependency on AI	Human-in-the-loop design ensures agents retain final decision authority
Auditability and compliance	Full logging of queries, retrieved clauses, model outputs, and human edits

These safeguards ensure the system improves operational efficiency **without increasing legal or compliance risk**.

Why a Pure “Buy” Approach Is Not Sufficient

While several enterprise AI copilots provide document search and generic RAG capabilities, they typically function as **general-purpose assistants** rather than workflow-aware decision tools.

The Contract & SLA Copilot requires several specialized capabilities that are difficult to achieve with an off-the-shelf solution alone:

- **Clause-level reasoning across multiple contract sections**
- **Ambiguity detection for conflicting amendments**
- **Mandatory human approval gates before customer commitments**
- **Audit logging aligned with service operations workflows**

Standard enterprise copilots often provide **document retrieval and summarization**, but they typically lack the ability to orchestrate the full **agentic workflow required for SLA interpretation and escalation handling**.

As a result, relying entirely on a vendor solution would likely require extensive customization while still limiting control over reasoning logic and escalation rules.

Therefore, a **hybrid architecture** — buying foundational infrastructure while building the specialized reasoning workflow — provides the best balance between **speed to MVP and long-term operational control**.

C4. Governance-Lite Checklist – Pilot Phase

Objective: Define the **minimum operational controls** required to safely launch a **4-week MVP pilot** for the Contract & SLA Copilot, ensuring:

- secure handling of **confidential contract data**
- **auditability of AI outputs**
- **human accountability in decision-making**
- measurable performance against **Golden Set benchmarks**

Governance Area	Control Specification & All Required Points	Acceptance & Monitoring Thresholds
Data & Access	Classification: Documents (MSAs/SLA Annexes) are Confidential ; Entitlement tables are Internal . Access: Restricted to 50 EMEA agents via SSO/RBAC; no anonymous use.	Audit Rule: 100% of sessions must be tied to a verified enterprise user ID for accountability.
Logging & Audit	Logging Fields: Capturing Query (Input), Clause IDs (Retrieval), AI Draft (Output), Agent Edits (HITL), and Latency (Performance) in the Centralized System Log .	Integrity: All logs must support the Golden Set evaluation and a full audit trail.
HITL Approvals	Gate 1: Service Agent reviews/verifies citations at Step 5 before sending. Gate 2: Contract Manager reviews flagged ambiguities at Step 6.	Human Accountability: No AI response is released to a customer without explicit human review/approval.
Evaluation Set	Golden Set: 50–100 labeled cases sourced from historical SLA tickets and validated by contract specialists. Include edge cases like conflicting amendments and ambiguous wording.	Accuracy Threshold: >=85% correct interpretation; >=90% citation grounding required to pass.
Escalation Plan	Monitoring: Weekly reviews of "Override Rate." Notification: Product Owner/Ops Lead notified if Accuracy <85% or Latency >60s.	Safety Pause: Pilot is paused if accuracy drops or "confidential" data leaks are detected.

D1. 30/60/90-Day Rollout Plan – Contract & SLA Answer Bot

Overview: Three-Phase Timeline

Days 1–30 PILOT	Days 31–60 EXPAND	Days 61–90 SCALE
Proof of concept with 10 Tier-1 EMEA service agents ; validate coverage/SLA/credits workflows , accuracy, and HITL usage (no data leaks)	Expand to all 50 EMEA service agents + contract managers ; fix top 3 issues (retrieval, latency, prompts); improve accuracy & resolution time	Full rollout as standard service desk workflow (SOP) ; ownership transferred to Service Ops; ROI readout vs €73,865 time savings target

Days 1–30: Pilot Scope (The First Steps)

Category	Specification & Strategy	Success Metrics & Guardrails
Users	10 Tier-1 EMEA Service Agents . Selected for high SLA query volume, mixed experience levels, and heavy contract ticket exposure to ensure representativeness for the full 50-agent rollout.	Selection Goal: Validate tool utility across different seniority levels before expanding to the "herd."
Workflow	Strictly Limited : (1) Coverage determination, (2) SLA interpretation, and (3) Credit eligibility. Excluded: Negotiations, multi-contract analysis, and complex legal escalations.	Scope Discipline: Prevents 4-week MVP from turning into an 18-month "slog".
Enablement	Day 1–2 Setup : 1-day training (RAG walkthrough, exercises, escalation rules, and RBAC governance). Artifacts: Prompt Library, Quick-Start Guide, and "Dos & Don'ts" for AI overrides.	Readiness: Ensures agents understand the Human-in-the-Loop (HITL) boundary before go-live.
Feedback Loop	Weekly Sprints : Review meetings with 10 agents, Product Owner, and Ops Lead. Focus: Human Override Rate, Accuracy audits, and Latency tracking. Measurement Source: System logs (usage, latency), audited sample set (50-case evaluation), and ticketing system data for handling time.	Iteration: Weekly ranking of top issues to drive prompt and retrieval refinements.
Success Gate	Go/No-Go at Day 30 : Pilot continues only if: (1) >=80% Adoption , (2) Zero Data Leaks , and (3) Stable HITL Usage (no blind acceptance). Measurement Method: Adoption via weekly active users (system logs); data leaks via audit logs; HITL stability via override rate tracking.	Pilot Minimum: Proves safety and utility before expanding to the remaining 40 agents.

Days 31–60: Expand (Finding the Rhythm)

Category	Specification & Strategy	Success Gates & Improvements
What Expands	User Growth: Onboarding the remaining 40 EMEA Service Agents and adding Contract Managers to operationalize the Step 6 escalation path.	Workflows: Full integration into ticket handling and the formal agent-to-specialist escalation chain.
System Integration	Lightweight Connectivity: CRM/ticketing context pre-fills and optimized document retrieval via the indexed contract repository to ensure the <60s latency target.	Technical Goal: Improve the "plumbing" of the RAG system based on real-world pilot data.
What Improves	Issue Resolution: Targeted fixes for the Top 3 Pilot Issues (e.g., retrieval gaps, latency, or ambiguity guardrails) .	Iterative Design: Deployment of Prompt Library v2 and enhanced citation formatting for better legal clarity.
Success Gate	Go/No-Go at Day 60: Transition to full scale only if: (1) Resolution Time <=10 minutes (vs. 37 min baseline). Measurement Method: Resolution time from CRM/ticketing timestamps; accuracy measured on a 50-case evaluation set and validated via human review sampling.	Quality Gates: (2) >=85% Answer Accuracy and (3) a Controlled Escalation Rate (specialists are not overwhelmed).

Days 61–90: Scale (Full Strength)

Category	Specification & Strategy	Success Metrics & Readout
Scale Plan	Standard Operating Procedure (SOP): The tool becomes the mandatory first step for all SLA queries, coverage checks, and credit determinations.	Workflow: Fully embedded into the daily service desk ticket resolution process .
Operating Model	Ownership: Service Operations Lead (Primary Owner) and Product Owner (Continuous Improvement). Support: Tiered structure from L1 (Agents) to L3 (Product/AI Team).	Cadence: Weekly KPI dashboard reviews and Monthly model performance audits (Accuracy, Overrides, Failure cases).
ROI Readout Plan	Day 90 Assessment: Formal evaluation comparing the 37-minute baseline against the ~6-minute target. Measurement Method: Time savings derived from system-recorded handling time reduction × loaded hourly cost (€55 baseline from Block B).	Measured Outcomes: Conversion of actual hours saved into Euro value, measured against the €73,865 time-savings target.

D2. KPIs – Definition, Data Source & Targets

11 KPIs across 4 categories — 3 Adoption (Leading), 3 Efficiency (Leading), 3 Quality (Leading & Lagging), 2 Business Outcomes (Lagging). Leading KPIs indicate early adoption and performance signals; lagging KPIs reflect realised business value and risk outcomes. STOP-AND-FIX triggers are defined where KPI deviations require pausing the rollout.

Category	KPI	Definition	Data Source	Baseline	Target (90 Days)
Adoption	Weekly Active Users (WAU)	% of total agents actively using the Copilot at least once per week	System usage logs (login + query activity)	0%	≥85% of 50 agents (~43 users)
Adoption	Feature Utilization Rate	% of eligible SLA/coverage/credit tickets where Copilot is used	CRM/ticketing system + AI usage logs	0%	≥70% of eligible tickets
Adoption	Avg. Queries per User	Average number of Copilot queries per agent per week	System logs (query count per user)	0	≥15 queries/user/week
Efficiency	Average Handling Time (AHT)	Average time to resolve SLA/contract queries	CRM/ticketing timestamps	37 minutes	≤6 minutes
Efficiency	System Latency (Time-to-Draft)	Time taken by Copilot to generate first response with citations	System logs (response time tracking)	N/A (manual process)	≤60 seconds
Efficiency	Agent Throughput	Number of SLA queries handled per agent per hour	CRM/ticketing system data	~1.6 queries/hour	≥6 queries/hour
Quality	Human Override Rate	% of AI responses modified or rejected by agents (HITL control)	System logs (AI output vs final response edits)	N/A	≤20%
Quality	Citation Grounding Rate	% of responses with correct clause-level citations	Evaluation set (50-case audit) + logging	N/A	≥90%
Quality	Escalation Accuracy	% of escalations correctly routed to contract managers (no unnecessary escalation)	CRM escalation logs + specialist review	N/A	≥85%
Business Outcome	Total Operational Value Captured	Value generated from time savings + error reduction vs baseline	CRM time data + cost model (€55/hr from	€0	≥€73,865 (time savings component)

	(€)		Block B)		
Business Outcome	Compliance Audit Pass Rate	% of responses meeting internal compliance/legal standards (no incorrect commitments)	Audit logs + compliance review samples	N/A	≥95%

D2A. Evaluation & Governance Plan (Pre-Launch + Safety Control)

1. Pre-Launch Evaluation Plan

1.1 Golden Test Set Design

Element	Specification
Test Set Size	100 cases
Source	Historical SLA/contract queries from CRM/ticketing system
Query Types	Coverage (40%), SLA (30%), Service Credits (30%)
Complexity Coverage	Includes edge cases (exceptions, amendments, conflicting clauses)
Purpose	Validate accuracy, grounding, and reasoning before pilot

1.2 Review Pack (Blind Evaluation Team)

Component	Description
Reviewers	Contract Managers (primary), Legal Analysts (validation)
Method	Blind comparison: AI answer vs Human expert answer
Evaluation Criteria	Correctness, completeness, citation validity
Process	Independent scoring → consolidated evaluation report

1.3 Success Benchmarks (Pre-Launch Gate)

Metric	Definition	Data Source	Threshold
Answer Accuracy	% of correct responses vs ground truth	Golden set evaluation	≥85%
Citation Grounding	% of responses with correct clause references	Clause validation	≥90%
System Latency	Time to generate response	System logs	≤60 seconds
Human Override Simulation	% of responses requiring edits	Test environment	≤20%
Condition: Pilot proceeds only if ALL thresholds are met			

2. Failure Modes & Protective Controls

2.1 Failure Mode Matrix

Failure Mode	Scenario	Risk Impact	Preventive Controls	Trigger (Stop Condition)
Hallucination	AI invents SLA terms or credits	Financial + legal exposure	Mandatory citation requirement; reject non-grounded answers; negative test cases	Citation grounding < 90%
Data Leak (RBAC Failure)	Access to wrong customer contract	Severe compliance breach	Role-Based Access Control; customer-level filtering; audit logs	Any high-severity data leak
User Rejection	Agents ignore or override AI outputs	Low adoption, no value realization	Training, prompt library, UI integration, feedback loops	Adoption < 60% OR override rate > 25%
Bias / Retrieval Skew	Model favors certain document types	Inconsistent decision quality	Balanced indexing; multi-document test cases	Accuracy drop in specific document types

3. Escalation & Kill-Switch Path

3.1 Human Intervention Flow

Step	Condition	Action	Owner
Step 5: Agent Review	Normal case	Agent reviews AI output	Tier-1 Agent
Step 6: Specialist Review	Low confidence / high-risk / unclear clause	Escalate to contract specialist	Contract Manager

3.2 Escalation Rules

Trigger Condition	Escalation Action
No clear clause found	Escalate to Contract Manager
Conflicting clauses	Specialist interpretation required
Customer-specific exceptions	Manual validation
High financial impact case	Mandatory escalation

3.3 Kill-Switch Authority

Role	Responsibility
Product Owner	AI system ownership; technical decision to pause
Service Operations Lead	Business ownership; operational decision to pause

3.4 Kill-Switch Conditions (Aligned with KPIs)

Condition	Threshold
Answer Accuracy Drop	<85%
Citation Grounding Failure	<90%
Human Override Rate Spike	>25%
Data Security Incident	Any high-severity leak
Compliance Failure	<95% audit pass rate

3.5 Recovery Protocol

Step	Action
1	Root cause analysis (retrieval, prompt, RBAC, etc.)
2	Fix implementation
3	Re-test using Golden Set
4	Approval from Product Owner + Ops Lead
5	Controlled restart of pilot

STOP-AND-FIX Triggers (Pause Rollout if Any of These Occur)

Trigger	Threshold	What It Indicates	Action (Stop & Fix)
Answer Accuracy Drop	< 85% (evaluation set)	Incorrect SLA/contract interpretations; model reliability issue	Pause rollout; Product Owner + Contract Manager review failure cases; fix retrieval/prompt logic; re-validate before restart
Human Override Rate Spike	> 25% (weekly)	Low trust in AI or poor output quality (HITL breakdown)	Analyze overridden queries; refine prompts/grounding; retrain agents if needed; resume after ≤20%
Citation Grounding Failure	< 90% grounded responses	Hallucinations or unsupported answers (high legal risk)	Pause customer-facing usage; fix retrieval/chunking; validate clause mapping; resume after ≥90%
Escalation Rate	> 30% (weekly)	Model lacks coverage OR agents	Review top escalations; improve missing clause coverage;

Spike		lack confidence	adjust confidence thresholds
System Latency Degradation	> 90 sec (sustained for 2 days)	Tool slows down workflow instead of improving it	Pause expansion; optimize retrieval/infrastructure; resume after ≤60 sec
Data Security Incident	Any HIGH-severity data leak	Breach of confidential data / RBAC failure	Immediate pilot suspension; security review within 24h; resume only after fix + approval
Compliance Failure	< 95% audit pass rate	Risk of incorrect legal commitments	Pause customer-facing outputs; legal review; tighten guardrails and approval flow
Adoption Failure	< 60% active users post-expansion	Low usability, trust, or training gap	Conduct agent interviews; improve UX/prompt library; reinforce training before scaling

D3. 10-Minute Pitch Script

SECTION 1 — The Strategic Problem

Today, answering a simple SLA or contract query takes approximately 37 minutes, based on a step-by-step process analysis. This delay is not caused by the complexity of the questions, but by a fragmented manual process where service agents must manually navigate Master Service Agreements, amendments, and appendices stored across disconnected systems. This manual lookup is a direct bottleneck to customer satisfaction and operational efficiency. More critically, it creates significant financial and legal exposure; every wrong interpretation leads to incorrect commitments, missed penalty credits, or compliance breaches. Currently, we are operating in a system where high-volume decisions are manual and knowledge is locked in unstructured documents. This is no longer scalable for an organization handling approximately 2,600 SLA queries annually—equivalent to 50 queries per week. And this gap is growing.

SECTION 2 — Why Now? The Agentic Shift

Traditional solutions like keyword search fail because contracts are defined by context, conditions, and exceptions, not just keywords. What has changed is the emergence of **Agentic RAG**. Unlike classic search, our Reasoning Agent evaluates natural language to understand how multiple clauses interact, resolving nuances that rule-based systems miss. This technology allows us to move from "searching documents" to "answering questions with evidence". By combining clause-level retrieval with natural-language reasoning, we can finally provide agents with grounded answers and verifiable citations. This enables us to automate the interpretation layer while maintaining a strict "legal ground truth".

SECTION 3 — MVP & Architecture Overview

The **Contract & SLA Copilot** is built on a **Pattern 4 Agentic Workflow** designed for high-stakes decision support. The architecture orchestrates four specialized agents: an **Extraction Agent** that interprets the customer intent; a **Retrieval Agent** that pulls precise clauses from the vector database; a **Reasoning Agent** that interprets conditions; and a **Response Agent** that drafts the final explanation. This is delivered through a lightweight UI panel embedded directly into the agent's workflow, providing cited answers and confidence indicators at a glance. Crucially, this is not a "black box" system. We enforce a mandatory **Human-in-the-Loop (HITL)** gate where agents must review, edit, or approve every suggestion before it reaches a customer. Every step—from raw input to the final human edit—is recorded in our **System Logging Layer** to ensure full auditability and continuous model improvement.

SECTION 4 — The Value Case & ROI

The financial impact of this deployment is measurable and immediate. Let’s quantify that. By reducing handle time from **37 minutes down to 6 minutes**, we capture a **31-minute saving per query**. Based on our baseline workload estimate of 50 queries per week—totaling 2,600 queries annually—this transition translates into over 1,000 hours of productive time freed for the department. We will validate this volume during the 30-day pilot by auditing historical service logs to ensure our ROI targets are fully realized. At a loaded cost of **€55 per hour**, this alone delivers **€73,865 in annual time savings**. When we add the mitigation of SLA penalties and the reduction in specialist rework, the total expected annual value reaches **€133,865**. This is a high-confidence ROI; the value is tied to specific KPIs—such as resolution time and override rates—that are tracked automatically through our system logs. We are essentially turning a cost center into an efficient, data-driven operation.

SECTION 5 — Rollout, KPIs, and Governance

Our deployment follows a disciplined **30/60/90-day rollout**. We begin with a **30-day pilot** featuring 10 Tier-1 agents, focusing strictly on coverage and credit queries. We will only expand if we meet our success gates of **≥ 80% adoption** and zero data leaks. By Day 60, we scale to the full 50-agent deployment, targeting an **85% answer accuracy** and **90% citation grounding**. Performance is managed via a dashboard of **11 KPIs** across adoption, efficiency, quality, and business outcomes. To ensure safety, we have implemented strict **"Stop-and-Fix" triggers**: if accuracy falls below 85% or any data exposure is detected, the pilot pauses immediately for specialist review. This governance-lite approach ensures we maintain human accountability and data security at every stage.

SECTION 6 — The Final "Ask"

To move forward, we are requesting an investment of **€38,000 to €55,000** to fund the 4-week MVP build and the initial pilot phase. This covers the infrastructure setup, document indexing, and the configuration of the agentic reasoning layer. In return, within 90 days, you will receive a validated system delivering a **€133,865 annual value**, with a projected **payback within the first year**. This project is not a broad, high-risk transformation; it is a targeted, controlled deployment with clear financial upside and a proven architecture. We are choosing to move from a slow, manual 37-minute process to a precise, 6-minute auditable system. We recommend immediate approval to begin the Day 1 pilot setup. This is not an AI experiment—this is a controlled operational upgrade with measurable ROI.